


Advances in DC deployment

- Conquering complexity.
 - Building racks of servers & complex cooling systems all separately is not efficient.
 - Package and deploy into bigger units:



Containers: Separating Concerns



Windows Azure Global Presence

North America Region

Europe Region

Asia Pacific Region



http://research.microsoft.com/en-us/um/redmond/events/cloudfutures2012/tuesday/Keynote_OpportunitiesAndChallenges_Yousef_Khalidi.pdf

Green Clouds

- Cloud Centers optimize life cycle costs and power use

$$\text{PUE} = \frac{\text{Total facility power}}{\text{IT equipment power}}$$

- <http://www.datacenterknowledge.com/archives/2011/05/10/uptime-institute-the-average-pue-is-1-8/>
- Average PUE = 1.8 (was nearer 3) ; Good Clouds are 1.1-1.2
- 4th generation data centers (from Microsoft) make everything modular so data centers can be built incrementally as in modern manufacturing
- <http://loosebolts.wordpress.com/2008/12/02/our-vision-for-generation-4-modular-data-centers-one-way-of-getting-it-just-right/>
- Extends container based third generation

Some Sizes in 2010

- <http://www.mediafire.com/file/zzqna34282frr2f/koomeydatacenterlectuse2011finalversion.pdf>
- 30 million servers worldwide
- Google had 900,000 servers (3% total world wide)
- Google total power ~200 Megawatts
 - < 1% of total power used in data centers (Google more efficient than average – **Clouds are Green!**)
 - ~ 0.01% of total power used on anything world wide
- Maybe total clouds are 20% total world server count (a growing fraction)

Some Sizes Cloud v HPC

- **Top Supercomputer** Sequoia Blue Gene Q at LLNL
 - 16.32 Petaflop/s on the Linpack benchmark using 98,304 CPU compute chips with 1.6 million processor cores and 1.6 Petabyte of memory in 96 racks covering an area of about 3,000 square feet
 - 7.9 Megawatts power
- **Largest (cloud) computing data centers**
 - 100,000 servers at ~200 watts per CPU chip
 - Up to 30 Megawatts power
 - Microsoft says upto million servers
- So **largest supercomputer** is around **1-2% performance of total cloud computing systems** with Google ~20% total

Data Center vs Supercomputers

- Scale
 - Blue Waters = 40K 8-core “servers”
 - Road Runner = 13K cell + 6K AMD servers
 - MS Chicago Data Center = 50 containers = 100K 8-core servers.


- Network Architecture

- Supercomputers: CLOS “Fat Tree” infiniband
 - Low latency – high bandwidth protocols
- Data Center: IP based
 - Optimized for Internet Access

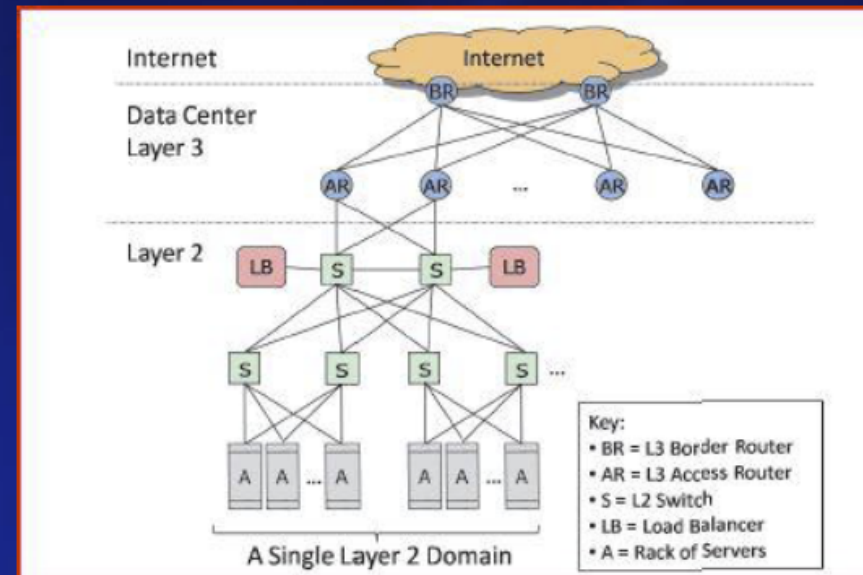
- Data Storage

- Supers: separate data farm
 - GPFS or other parallel file system
- DCs: use disk on node + memcache

Fat tree network



Standard Data Center Network



Cloud Industry Players

Players: Providers

Programmatic access via Web Services and/or Web APIs

“Pure” virtualized resources

CPU, memory, storage, and bandwidth

Data store



versus



Virtualized resources plus application framework
(e.g., RoR, Python, .NET)

Imposes an application and data architecture

Constrains how application is built



Google App Engine



[MM]

Players: Cloud Intermediaires

Resells (aspects of) raw cloud resources, with added value propositions

- Packaging resources as bundles

- Facilitating cloud resource management,
e.g., setup, updates, backup, load balancing, etc.

- Providing tools and dashboards

Enabler of the cloud ecosystem



[MM]







Players: Application Providers

Software as a Service (SaaS):

Applications provided and consumed over the Web

Infrastructure usage (mostly) hidden



-  **Gmail**
Email with up to 25 GB of storage per custom email address, mail search tools and integrated chat.
-  **Google Docs**
Create, share and collaborate on documents in real-time.
-  **Google Calendar**
Coordinate meetings and company events with sharable calendars.
-  **Google Sites**
One-stop sharing for team information.
-  **Google Talk**
Free text and voice calling around the world.
-  **Security and compliance**
Set email policies and recover deleted messages.

Cloud Talk

▲ *The Public Cloud “Big Four”*


▲ *Amazon*

▲ *Google*

▲ *Microsoft*

▲ *Salesforce.com*

johnmwillis.com



Cloud Talk

▲ *The Private Cloud “Big Four”*

▲ *IBM*

▲ *VMware*

▲ *Sun/Oracle*

▲ *3Tera*

johnmwillis.com

